


Linear Regression I

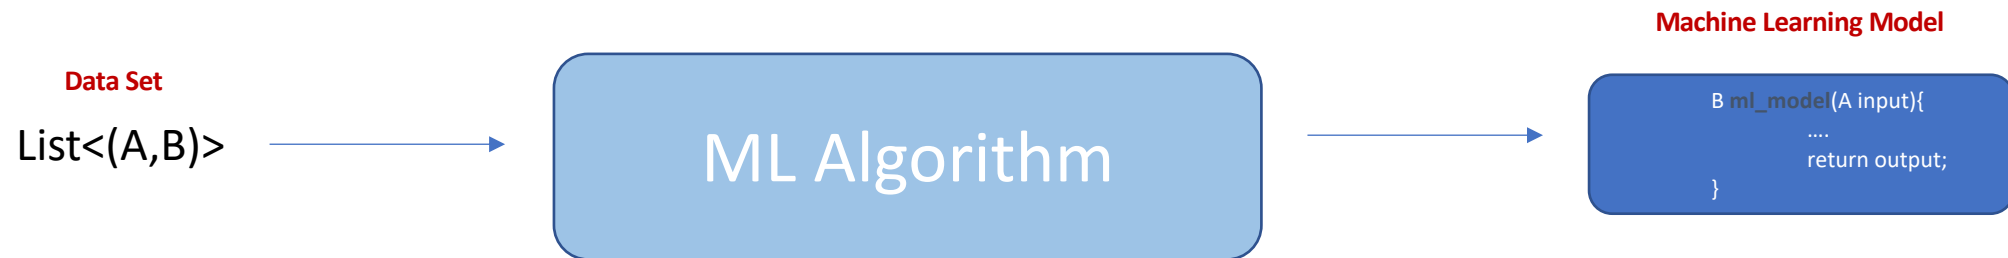
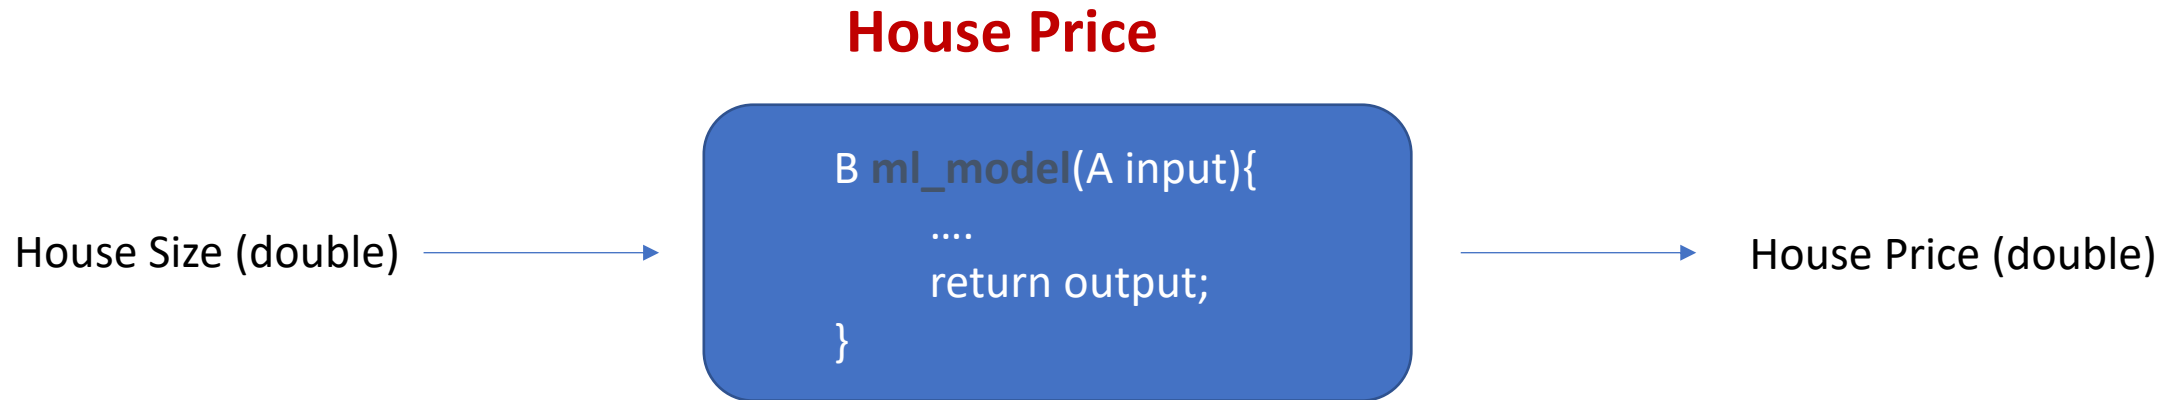
Andrew Ng's Machine Learning Course



Machine Learning



Predict House Price given House Size



Predict House Price given House Size

Data table


size in feet ²	price in \$1000's
2104	400
1416	232
1534	315
852	178
...	...
3210	870

Predict House Price given House Size

Data table

size in feet ²	price in \$1000's
2104	400
1416	232
1534	315
852	178
...	...
3210	870

House sizes and prices



Terminology and Notation

Terminology

Training set: Data used to train the model

size in feet ²	price in \$1000's
2104	400
1416	232
1534	315
852	178
...	...
3210	870

Notation:

x = "input" variable
feature

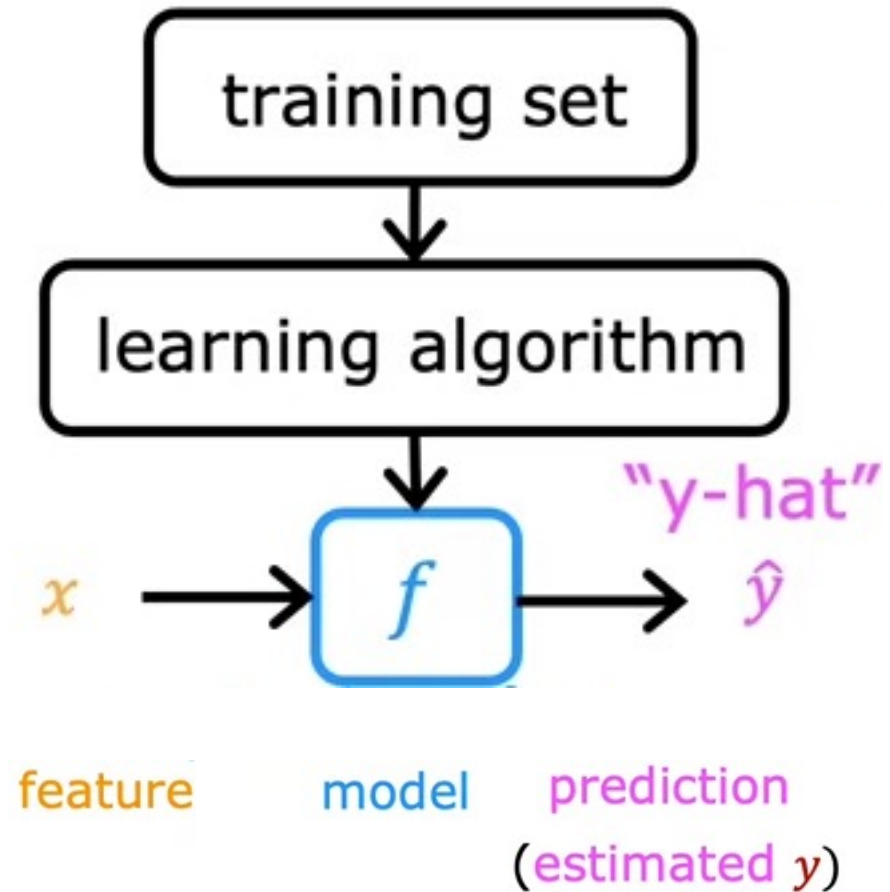
y = "output" variable
"target" variable

m = number of training examples

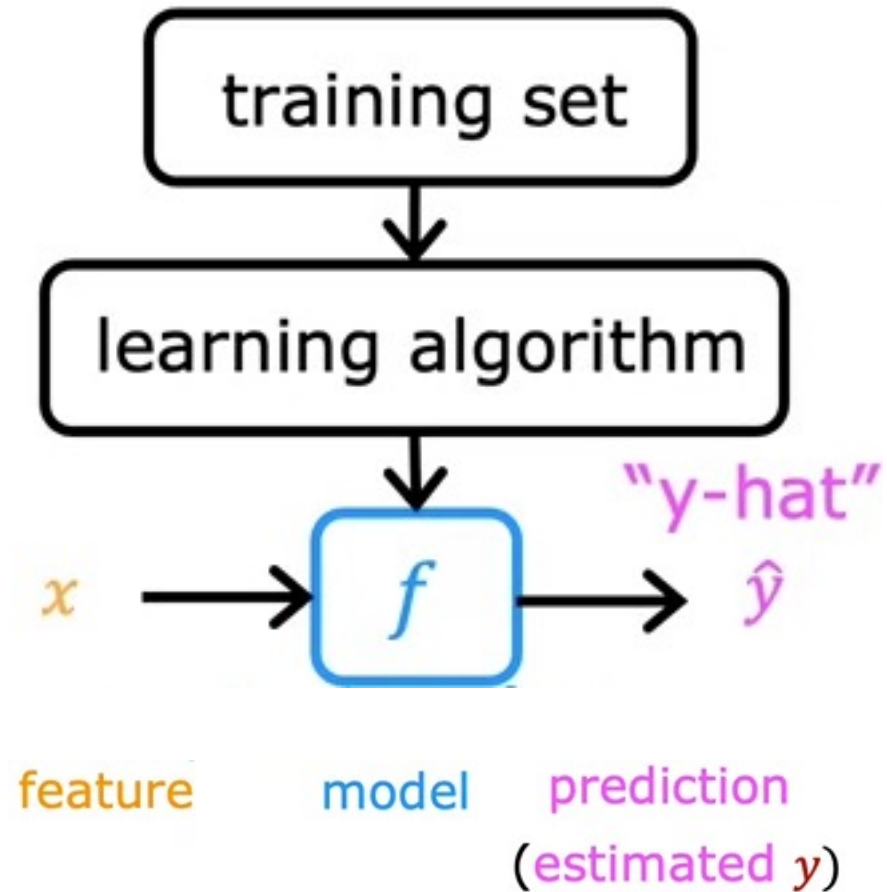
= single training example

$(x^{(i)}, y^{(i)})$ = i^{th} training example


Machine Learning Model



Machine Learning Model



How to represent f ?



Linear Regression

Training set

size in feet ² (x)	price \$1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

Model: $f_{w,b}(x) = wx + b$




w, b : parameters

What do w, b do?

Linear Regression

$$\underline{f_{w,b}}(x) = wx + b$$

$f(x)$



Exercise 1: Model Function

Exercise 1


Complete the function `compute_model_output` to compute the output of a linear regression model given a vector of x values and a two parameters w and b .

Note: The argument description `(ndarray (m,))` describes a Numpy n-dimensional array of shape $(m,)$. `(scalar)` describes an argument without dimensions, just a magnitude.

Note: `np.zeros(n)` will return a one-dimensional numpy array with n entries

```
[ ] def compute_model_output(x, w, b):  
    """  
    Computes the prediction of a linear model  
    Args:  
        x (ndarray (m,)): Data, m examples  
        w,b (scalar)      : model parameters  
    Returns  
        y (ndarray (m,)): target values  
    """  
    m = x.shape[0]  
    f_wb = np.zeros(m)  
    for i in range(m):  
        f_wb[i] = 0# Introduce code here  
  
    return f_wb
```

Cost Function



$$\hat{y}^{(i)} = f_{w,b}(x^{(i)})$$

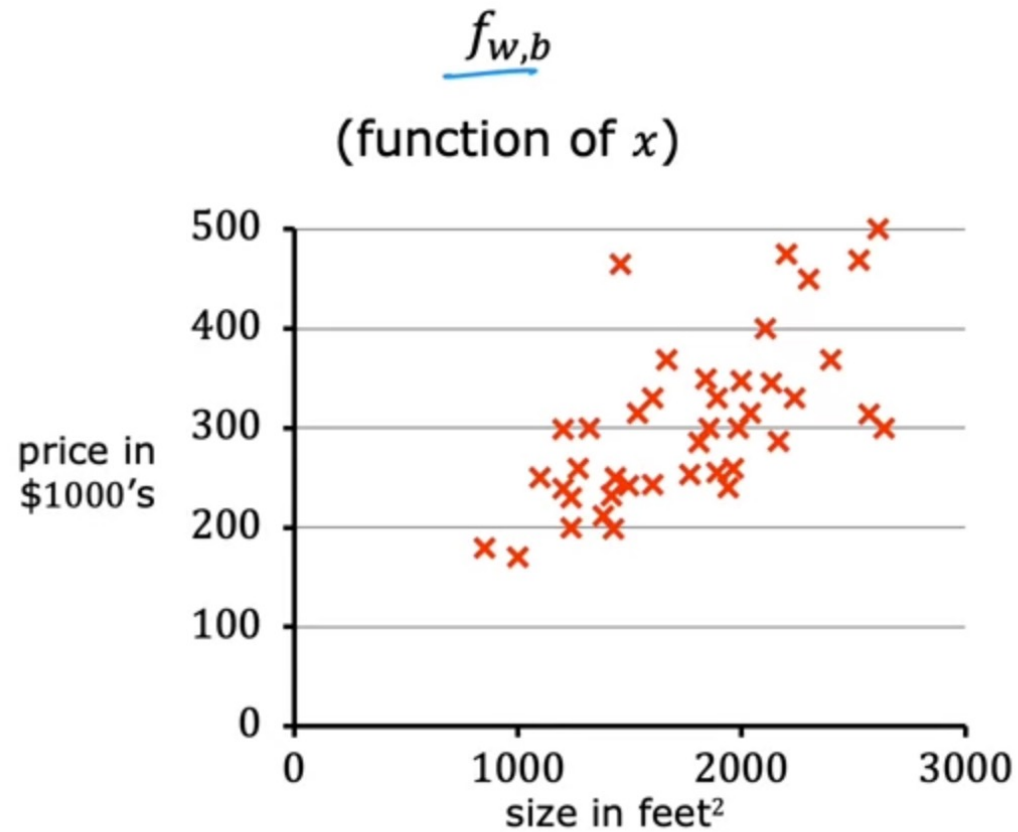
$$f_{w,b}(x^{(i)}) = wx^{(i)} + b$$

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Find w, b :

$\hat{y}^{(i)}$ is close to $y^{(i)}$ for all $(x^{(i)}, y^{(i)})$.

Cost Function




J
(function of w, b)

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Objective

minimize $J(w, b)$
 w, b

Cost Function for Linear Regression



Exercise 2: Cost Function

▼ Exercise 2: Cost Function

Complete the code of the `compute_cost` method below to:

- Iterate over the training examples, and for each example, compute:
 - The prediction of the model for that example

$$f_{wb}(x^{(i)}) = wx^{(i)} + b$$

- The cost for that example

$$cost^{(i)} = (f_{wb}(x^{(i)}) - y^{(i)})^2$$

- Return the total cost over all examples

$$J(w, b) = \frac{1}{2m} \sum_{i=0}^{m-1} cost^{(i)}$$

- Here, m is the number of training examples and \sum is the summation operator

If you get stuck, you can check out the hints presented after the cell below to help you with the implementation.

```
def compute_cost(x_train, y_train, w, b):  
    """  
    Computes the cost function for linear regression.  
  
    Args:  
        x_train (ndarray): Shape (m,) Input to the model (Population of cities)  
        y_train (ndarray): Shape (m,) Label (Actual profits for the cities)  
        w, b (scalar): Parameters of the model  
  
    Returns  
        total_cost (float): The cost of using w,b as the parameters for linear regression  
        to fit the data points in x and y  
    """  
    # number of training examples  
    m = y_train.shape[0]  
  
    total_cost = 0  
    for i in range(m):  
        f_i = 0  
        cost_i = 0 # Hint: Use y_train[i]  
        total_cost += cost_i  
  
    total_cost/=2*m  
  
    return total_cost
```

► [Click for hints](#)

Minimizing the Cost Function

Have some function $J(w, b)$

Want $\min_{w, b} J(w, b)$


Outline:

- Start with some w, b

- Keep changing w, b to reduce $J(w, b)$

- Until we settle at or near a minimum

Minimizing the Cost Function



Gradient Descent Algorithm

repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

Gradient Descent Algorithm

repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$


$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

$$J(w)$$

$$w = w - \alpha \frac{\partial}{\partial w} J(w)$$


$$\min_w J(w)$$

GD: Moving along the gradient




$$J(w)$$
$$w = w - \alpha \frac{\partial}{\partial w} J(w)$$

$$\min_w J(w)$$




GD: Learning Rate



$$J(w)$$
$$w = w - \alpha \frac{\partial}{\partial w} J(w)$$

$$\min_w J(w)$$




GD: Convergence


$$J(w)$$

$$w = w - \alpha \frac{\partial}{\partial w} J(w)$$

$$\min_w J(w)$$



GD: Local Minima



$$w = w - \alpha \frac{d}{dw} J(w)$$

GD for Linear Regression

Linear regression model

$$f_{w,b}(x) = wx + b$$

Cost function

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Gradient descent algorithm

repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

}

GD for Linear Regression

Linear regression model

$$f_{w,b}(x) = wx + b$$

Cost function

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

Gradient descent algorithm

repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

}


repeat until convergence {

$$w = w - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

}

GD for Linear Regression



Exercises 3 and 4: Gradient Descent

Exercise 3

Please complete the `compute_gradient` function to:

- Iterate over the training examples, and for each example, compute:

- The prediction of the model for that example

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b$$

- The gradient for the parameters w , b from that example

$$\frac{\partial J(w, b)}{\partial b}^{(i)} = (f_{w,b}(x^{(i)}) - y^{(i)})$$

$$\frac{\partial J(w, b)}{\partial w}^{(i)} = (f_{w,b}(x^{(i)}) - y^{(i)})x^{(i)}$$

- Return the total gradient update from all the examples

$$\frac{\partial J(w, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} \frac{\partial J(w, b)}{\partial b}^{(i)}$$

$$\frac{\partial J(w, b)}{\partial w} = \frac{1}{m} \sum_{i=0}^{m-1} \frac{\partial J(w, b)}{\partial w}^{(i)}$$

- Here, m is the number of training examples and \sum is the summation operator

If you get stuck, you can check out the hints presented after the cell below to help you with the implementation.

```
# UNQ_C2
# GRADED FUNCTION: compute_gradient
def compute_gradient(x, y, w, b):
    """
    Computes the gradient for linear regression
    Args:
        x (ndarray): Shape (m,) Input to the model (Population of cities)
        y (ndarray): Shape (m,) Label (Actual profits for the cities)
        w, b (scalar): Parameters of the model
    Returns
        dj_dw (scalar): The gradient of the cost w.r.t. the parameters w
        dj_db (scalar): The gradient of the cost w.r.t. the parameter b
    """

    # Number of training examples
    m = x.shape[0]

    # You need to return the following variables correctly
    dj_dw = 0
    dj_db = 0

    ### START CODE HERE ###

    ### END CODE HERE ###

    return dj_dw, dj_db
```

► Click for hints

Exercise 4

You will now find the optimal parameters of a linear regression model by using batch gradient descent. Recall batch refers to running all the examples in one iteration.

- TASK:** Implement the updating equation of gradient descent

$$\begin{aligned} b &:= b - \alpha \frac{\partial J(w, b)}{\partial b} \\ w &:= w - \alpha \frac{\partial J(w, b)}{\partial w} \end{aligned} \quad (1)$$

- A good way to verify that gradient descent is working correctly is to look at the value of $J(w, b)$ and check that it is decreasing with each step.
- Assuming you have implemented the gradient and computed the cost correctly and you have an appropriate value for the learning rate α , $J(w, b)$ should never increase and should converge to a steady value by the end of the algorithm.

```
def gradient_descent(x, y, w_in, b_in, cost_function, gradient_function, alpha, num_iters):
    """
    Performs batch gradient descent to learn theta. Updates theta by taking
    num_iters gradient steps with learning rate alpha

    Args:
        x : (ndarray): Shape (m,)
        y : (ndarray): Shape (m,)
        w_in, b_in : (scalar) Initial values of parameters of the model
        cost_function: function to compute cost
        gradient_function: function to compute the gradient
        alpha : (float) Learning rate
        num_iters : (int) number of iterations to run gradient descent

    Returns
        w : (ndarray): Shape (1,) Updated values of parameters of the model after
            running gradient descent
        b : (scalar) Updated value of parameter of the model after
            running gradient descent
    """
```

Exercise 5: Vectorization using Numpy

```
def compute_model_output(x, w, b):  
    """  
    Computes the prediction of a linear model  
    Args:  
        x (ndarray (m,)): Data, m examples  
        w,b (scalar)      : model parameters  
    Returns  
        y (ndarray (m,)): target values  
    """  
  
    """  
    m = x.shape[0]  
    f_wb = np.zeros(m)  
    for i in range(m):  
        f_wb[i] = x[i]*w + b  
    """  
  
    f_wb = x*w + b  
  
    return f_wb
```